

Inside-Out Code Auto-Completion

Stanford CS224N Custom Project
Mentor: Mina Lee
Grading option: 3 (graded)

Gabriel Poesia
Department of Computer Science
Stanford University
poesia@stanford.edu

Lauren Gillespie
Department of Computer Science
Stanford University
gillespl@stanford.edu

Scott Viteri
Department of Computer Science
Stanford University
sviteri@stanford.edu

Abstract

Auto-complete has become a ubiquitous feature in modern-day typing systems. Traditional systems typically try to predict the next characters or words the user is about to type. In contrast, keyword-based auto-completion systems expect the user to type keywords, then fill in the least informative parts of the sentence, and have been successful for a variety of natural language tasks. In this project, we implement a keyword-based auto-complete technique for programming languages, as auto-complete systems in modern IDEs are still mostly left-to-right. We use three programming languages with different verbosity levels: Python, Haskell and Java. We evaluate different encoding schemes paired with a neural decoder with various trade-offs in translation accuracy and compression, and evaluate their accuracy and robustness.

1 Introduction

Auto-complete has become a ubiquitous feature in almost all modern-day typing systems. However, almost all of these auto-complete systems utilize left-to-right (LTR) predictive auto-completion. Yet the most important information in a sentence or phrase isn't guaranteed to be clustered at the beginning. A more efficient system instead would enable users to convey the most important information of the sentence or phrase in the form of keywords, so instead of predicting the rest of the sentence from left-to-right, the system predicts the full sentence from its keywords, in an *inside-out* fashion.

For example, in the LTR setting, if a user wants to communicate the sentence “I want to see Lion King at 10 o'clock pm” she will first type “I want to see”. Even then, it will be difficult for the system to infer the rest of the sentence, since the first few words have such low information density. In the keyword-based setting the user might only have to type “Lion King 10 pm”.

Similar bottlenecks occur when programming. For example, to when typing `for i in range(10):` Pythia¹ only auto-completes `for i in r to for i in range`, and `(10)` is manually typed. However, typing `for i in 10` would convey the same information to a human user in fewer characters. A keyword-based system might be able to auto-complete this shorthand.

In this paper, we describe and analyze a character-level auto-complete system for lines of source code. We experiment with several choices that place an auto-completion scheme in the compactness-

¹a left-to-right auto-complete system

accuracy spectrum. In Section 2, we survey related work. Section 3 formalizes the auto-completion problem and explains our solution. Section 4 describes our experimental setup. Section 5 discusses the results and examples from the various proposed auto-completion schemes. Finally, Section 6 summarizes our results.

2 Related Work

Previous work implemented a neural-based auto-complete system for natural language tasks[?]. This system implements a context-sensitive neural encoding scheme which compresses target sequences. While Lee et. al. implements inside-out auto-completion for natural language tasks, our system specifically focuses on programming language implementation of inside-out auto-completion.

Little et. al. also proposed a keyword-based Java auto-complete system[?]. However, their system only generates method calls, variable references, and constructor calls, while ours generates arbitrary lines of code. Their implementation operates on a word level, while ours is character-based. Also, their system is not robust to abbreviations of keywords themselves; the system cannot match keywords that are not character-for-character present in the table of keywords translations. Finally, their system is language-specific while our proposed method is language-agnostic.

Reiss 2009 implements a keyword-based function generator for Java code [?]. In this system keywords are used as search keys into a database of existing programs, and matched programs undergo a series of program transformations. These program transformation must be written separately for each language. The main purpose of this program is to search for pre-existing programs, rather than generate wholly new ones.

3 Approach

We now describe our approach in building a code auto-complete system. We define an *auto-completion scheme* A to be a pair of functions $A = (E, D)$, both mapping strings to strings. Suppose the user wants to type a certain string s . By applying the *encoder* of A , we get $s' = E(s)$, a string that the user can type instead of the potentially longer s , and use the decoder function $D(s')$ to get back the original string s .

Two features would be desired of a scheme A . First, we would like the encoded strings to be small, as the overall goal is to save the user typing effort. More precisely, we want $\mathbb{E}_{s \in C} |E(s)|$ to be as small as possible, where the expectation is taken over the probability distribution of lines of code in a certain programming language. Furthermore, we want the decoder to be able to precisely determine the original line of code after it has been encoded, i.e. we want $\mathbb{E}_{s \in C} (\mathbb{P}(s = D(E(s))))$ to be as close to 1 as possible. These two goals are clearly conflicting: the most precise encoding scheme would be the identity scheme, which is not compact, while any scheme that shortens the input will introduce uncertainty in the decoding. We therefore study several auto-completion schemes that fall into different points of this trade-off.

Finally, we also want an auto-complete scheme to be intuitive to users. A compact and precise scheme could be built by assigning common lines of code to short sequences of characters that rarely occur in real code (or even do not belong to the programming language at hand). To keep the encoded strings related to the user’s intent, as in [?] we restrict ourselves to encoding schemes that output a *sub-sequence* of the input characters. In other words, the encoder can only drop, but not add, characters to the input.

3.1 Encoder

The *encoder* side of an auto-complete scheme aims at dropping the least informative characters in such a way that the decoder is still able to recover the input. We consider two kinds of encoders: *ad hoc* encoders, which are plain algorithms that a neural decoder is trained to invert, and a *neural context-sensitive* encoder that is trained end-to-end alongside the decoder. We designed the following *ad hoc* encoders:

Rules. Many programming languages have common commands that are long to type. For example, to write a line to the standard output in Java, one would use the `System.out.println()`

method (or a similar alternative). From a list of such common commands and keywords, several modern Integrated Development Environments (IDEs) have a table of intuitive abbreviations that expand to the long commands. Some Java IDEs, for instance, expand `sout` to `System.out.println`, which users can easily memorize. Our *Rules* encoder applies a manually-crafted table of abbreviations to the input lines of code. Examples of abbreviations we have for Python are `T` \rightarrow `True` and `rev` \rightarrow `reverse`.

Uniform. While *Rules* is very intuitive, it is also very rigid, as all abbreviations are defined *a priori*. On the other extreme, we can think of simply dropping characters at random. This is the *Uniform* encoder, which is parameterized by the probability p of dropping each character, independently. Therefore, *Uniform* can abbreviate any string s into a string that has expected length $p \times |s|$.

Frequency. The flexibility of *Uniform* also has a cost: while it’s easy to identify common sequences of characters like `System.out.println` even when any 30% of the characters are removed, rarer sequences are hardly identifiable, as there are few examples of them. The same happens for short sequences (like operators) that might be entirely erased during encoding. As a balance between the focus of *Rules* on frequently-used keywords and the breadth of *Uniform*, the *Frequency* encoder first ranks character n -grams by their frequency in the training set, and then replaces n -grams by a concatenation of their first and last characters until the length of the input is reduced to a desired fraction p . For example, if $p = 0.8$, the length of the input is reduced by 20%. After some experimentation, we decided to use $n = 5$ as our n -gram size. One example of a rule that the *Frequency* encoder uses is `print` \rightarrow `pt`.

Noisy Frequency. While a system trained with *Frequency* achieves a good compression rate and accuracy, in practice we observed this encoder is not robust to the user using different abbreviations than the first and last characters of the n -gram (we show this analysis in Section 4). This is important for use in the real world, as we also want users not to have to need much adaptation to use the system. With this motivation, we devised the *Noisy Frequency* encoder, which is also parameterized by the desired fraction p of characters to be kept. The goal is to train *Frequency* to be also robust to noise. Starting from the input, in each step this encoder either removes one character at random or applies the same rule as *Frequency* that abbreviates one n -gram. Both options are equally likely and subsequent decisions are independent.

Besides these *ad hoc* encoders, we also implemented a *neural context-sensitive* encoder that is trained end-to-end with the decoder. We use the same model as Lee et al. [?], where an LSTM model outputs a probability of dropping each character in the input. We then sample from that probability distribution to get the encoded string, which is then fed into the decoder that tries to reconstruct the original output. Unfortunately, despite our best efforts, we did not manage to get useful results out of this model. While we used their same constrained objective function, we still got the results that are reported for the less stable linear objective, where the encoder degenerates during training to either keep all characters (minimizing the decoder loss) or drop all characters (minimizing the average length of the encoded strings). We suspect that, since our system is character-level, which makes input sequences significantly larger than the ones this encoder architecture was originally designed for, the model might take more time to leave trivial local minima than what we had available. Therefore, all results and analyses we show only consider *ad hoc* encoders.

3.2 Decoder

We pair all described encoders with a neural decoder implemented as a standard sequence-to-sequence architecture with multiplicative attention, where the input (the abbreviated string) is first encoded by a bidirectional LSTM, and then decoded into the expanded string by a uni-directional LSTM. We jointly train character embeddings with our models, which was observed to be better than using one-hot encodings.

3.2.1 Attention-based copy mechanism

As in [?], we also integrated a *copy mechanism* in our model, as described in [?]. We observed that the model begins training with a smaller loss than without copy, but the final accuracy is decreased by 10% on average in all settings. After analyzing running examples, we observed that the attention

distribution is not precise enough in the case of a character-based code model to inform the characters to be written to the output.

We notice that there are significant differences between our use case and the scenario that attention-based copy was designed for. First, the sequences of characters we deal with are on average much longer than the number of tokens in a typical sentence in natural language. More precisely, our average line of Python code had 36.5 characters, whereas the auto-complete system in [?] was only tested in sequences of at most 16 tokens. Also, the alphabet size in our case is 128, from ASCII encoding: much smaller than the tens of thousands of words that are typically used in the vocabulary of natural language systems. These differences likely make attention-based copy mechanisms inappropriate to our case, and therefore we did not use it in our final model.

3.2.2 *Hard-copy mechanism*

Motivated by the fact that the output of the decoder is a super-sequence of its input, we designed a custom *hard-copy* mechanism that is not attention-based. If the original input had n characters and the encoded input has $n' < n$ characters, we know at training time that the decoder has to copy all n' characters from its input at some point, and also insert $n' - n$ characters in between. Furthermore, we know that the copied characters will come in order, and because we know which characters were removed, we can also compute the expected sequences of actions of the decoder, being either to insert a character or copy the next character from the input. For example, if the string `print` was shortened to `pt`, we expect the decoder to first copy `p`, then do three insertions of `r`, `i` and `n` and finally copy `t`.

We experimented with augmenting the decoder output with a special character that triggers a copy of the next character from the input that has not yet been copied. In this way, the copy is *non-parametric*, i.e. the model doesn't need to specify which character to copy. We observed that, at first, if the minority of the characters are dropped by the encoder, most of the expected output characters are the special "copy" token, and the decoder quickly learns to always copy at every step. Under this framework, the model can achieve zero loss on all characters that have not been dropped if it outputs a probability of 1 to the copy character at all time steps. Thus, it gets stuck at this local optimum of only copying the compressed representation, which clearly fails in all test examples where at least one character was removed. If we drop more than the majority of the characters, the model does not learn to output "copy" at all time steps, but also doesn't learn useful associations as the compressed representation simply contains too little information to predict the original line. We therefore ended up not using this *hard-copy* mechanism in our decoder.

4 Experiments

We evaluated all *ad hoc* encoders described in Section 3, paired with the neural decoder. Our investigation is guided by three research questions:

- RQ1:** How do the different auto-complete schemes compare in accuracy and robustness?
- RQ2:** What is the impact of the compression rate of an auto-complete scheme in its accuracy?
- RQ3:** How do different programming languages impact the accuracy of auto-complete systems?

4.1 Data

We built a dataset comprised of code extracted from public GitHub repositories written in Python, Haskell, and Java. We ignored comments, indentation and lines that are either fewer than 10 or greater than 80 characters long. In total, for each programming language, we used 600k unique lines of code for training, 80k for validation and hyper-parameter tuning, and 80k for testing.

4.2 Experimental Details

All experiments were executed on a Intel(R) Xeon(R) E5-2690 machine running Ubuntu Server 16.04, with 56GB of RAM, and a single NVIDIA Tesla M60 GPU with 8GB of memory.

For the neural decoder, we used a 512-dimensional single-layer LSTM model. We jointly trained 50-dimensional character embeddings with each encoder-decoder pair. Each model was trained for 4h using the Adam optimizer, with a learning rate of 0.0001, $\beta_1 = 0.9$ and $\beta_2 = 0.999$.

	Rules Dec.	Uniform Dec.	Frequency Dec.	NoisyFrequency Dec.
Rules Enc.	0.943 / 0.980	0.643 / 0.805	0.210 / 0.385	0.182 / 0.450
Uniform Enc.	0.008 / 0.015	0.181 / 0.340	0.002 / 0.010	0.070 / 0.250
Frequency Enc.	0.000 / 0.005	0.077 / 0.270	0.592 / 0.765	0.037 / 0.115
NoisyFrequency Enc.	0.000 / 0.000	0.020 / 0.060	0.000 / 0.000	0.248 / 0.410

Table 1: Top-1 Accuracies / Top-5 Accuracies for Encoder/Decoder Pairs.

We trained all 4 *ad hoc* encoders described in Section 3. For *Uniform*, *Frequency* and *NoisyFrequency*, we set their parameter p to 0.8 unless noted otherwise, which means *Uniform* drops each character independently with probability 0.2, and *Frequency* and *NoisyFrequency* remove characters until the input is reduced by at least 20%. Finally, *Frequency* and *NoisyFrequency* are always set to operate on character 5-grams.

4.3 Accuracy and Robustness of Auto-Complete Schemes

To compare the different proposed auto-complete schemes, we first trained neural decoders for each all 4 encoders: *Rules*, *Uniform*, *Frequency* and *NoisyFrequency*. Then, we split each scheme into an encoder-decoder pair, and evaluate all decoders when running on the output of all encoders. Ideally, an auto-complete scheme should be accurate when its encoder and decoder are paired, but we’d also like the decoder to be robust and be able to correctly reconstruct input abbreviated by different encoders, as we cannot expect all users to use the same abbreviations.

Table 1 shows the accuracies of both top-1 and top-5 predictions for all encoder/decoder pairs. We note that *Rules* is very accurate on itself, having a 98% top-5 accuracy. However, a decoder only trained along the *Rules* encoder is not robust: its top-5 accuracy is at most 2% for other encoders. *Frequency* decoder has a high accuracy when the input comes from the corresponding encoder (76.5% top-5 accuracy), but it is not robust: its top-1 accuracy drops to zero on both *Uniform* and *NoisyFrequency*. This indicates that it cannot perform expansions that are not in the computed table of frequent n -grams.

As expected, *NoisyFrequency* is slightly more robust, having its worst top-5 accuracy of 11.5% on the *Frequency* encoder. Surprisingly, its accuracy increases when paired with the *Uniform* encoder, going up to 25%, suggesting that the noise makes the model not specialize to frequent n -grams necessarily. Finally, *Uniform* decoder has the best performance on the *Rules* encoder besides its across all decoders other than *Rules* itself. This suggests that *Uniform* successfully learns to pick up common-sense patterns, even though it is trained on purely random abbreviations.

These results indicate a positive correlation between the regularity of an encoder and the accuracy of its decoder, but this accuracy trades-off with robustness. *Rules* is the most predictable encoder, and therefore leads to the most accurate auto-complete scheme; however, it is the least robust. On the other extreme, *Uniform* has the worst result on itself, but was shown to be more robust than other schemes, and is able to decode *Rules* fairly accurately even without specialized training. These observations answer our first research question.

4.4 Impact of compactness on accuracy

We now aim at understanding the impact of the compactness of an auto-complete scheme on its accuracy. To that end, we trained *Uniform* decoders with four different probabilities of keeping each character: 0.6, 0.7, 0.8 and 0.9. Figure 1 shows the top-1 accuracies of each of the obtained decoders when expanding input that was compressed to each of these sizes.

First, we note that the learned *Uniform* decoders do not overfit the relationship between the length of the input and the length of the output, which has a fixed average during training. For example, for *Uniform 0.8* decoder, if the input has 8 characters, it could potentially expect the output to have 10 characters based on its training procedure. However, in practice, this does not happen. As one would desire, the performance of all decoders increase monotonically when we drop less characters. The worst performance of all decoders happen when the input has only 60% of the original characters. When 90% of the input is kept, the decoders perform at their best.




Figure 1: Top-1 accuracy of *Uniform* decoders trained with different compression rates.

Also, we notice an interesting relationship between robustness of the learned scheme and the amount of noise present in the input. All decoders have a higher accuracy than *Uniform 0.9* when 90% of the characters are kept, even if *Uniform 0.9* decoder was specialized for this case. This suggests that seeing more noise in the input helps other decoders in easier cases as well. On the other extreme, *Uniform 0.6* performs worse than *Uniform 0.7* and *Uniform 0.8* when 90% of the characters are kept on average. As much as noise helps in robustness, too much noise can also hinder the decoder from learning useful representations. Most likely, *Uniform 0.7* and 0.8 are closer to the sweet spot in the amount of noise during training, since their performance was mostly equivalent on all noise scenarios and superior to the other two decoders. This result provides insight to answer our second research question.

4.5 Impact of different programming languages on auto-complete schemes

Finally, we evaluate how different programming languages impact the performance of auto-complete schemes. First, we notice that the average line length of a programming language determines how many characters are dropped on average by *Uniform*, which alone makes it easier or harder to recover all dropped characters accurately. Among the 3 languages we analyze, Python has the smallest average line length, at 35.4 characters. Haskell is a close second, with 35.9, while Java has the longest average lines, with 39.3 characters. To let this difference aside, in this experiment we instead drop a constant number of 5 characters of each line, instead of a variable fraction of the input.

Table 2 shows top-1 accuracy results for all 3 programming languages, and also for different line lengths. Of these 3 languages, Haskell is commonly known as the most concise, and Java as the most verbose. Our results match this common perception: even though line lengths for Haskell and Python are very similar on average, it is significantly harder to auto-complete Haskell code. When analyzing individual lines, we see a larger prevalence of one-character names in that language, for example. When such a character is dropped, it is not reasonable to expect it to be recovered by the decoder. This makes longer lines harder in Haskell than shorter lines, because there are more opportunities for making this kind of mistake. The same phenomenon is observed in Python. In Java, however, longer lines are easier for the decoder than shorter lines. We noticed that the longest Java lines usually contain common long class names, such as `ConcurrentHashMap` or `BufferedOutputStream`, which tend to appear twice in the same line in variable declarations. Therefore, longer lines in Java tend to have more redundancy than longer lines in Python or Haskell, which are usually very dense. These results let us conclude that it is significantly easier to build accurate auto-complete systems for more verbose programming languages. This answers our third research question.

	Python	Java	Haskell
All lines	0.201	0.353	0.154
Short lines	0.182	0.239	0.124
Long lines	0.174	0.373	0.120

Table 2: Top-1 Accuracies for *Uniform* trained with constantly dropping 5 characters for Python, Java, and Haskell. Long lines contain more than 65 characters and short lines contain less than 25 characters.

5 Analysis

Suppose the user wants to type `for i in range(10):`. The rules-based decoder correctly translates `for i in rng(10):` to `for i in range(10):`, but incorrectly translates `for i in rg(10):` to `for i in rg(10):`. Rules-based works for the abbreviation 'rng', because it was explicitly in the translation list. But if the abbreviation is not on the list, the decoder struggles: the rules-based decoder will only work on a small range on abbreviations. We want our model to be flexible with respect to the various encodings that a programmer might use.

One way to ensure flexibility is to separately decide whether to drop each character with some fixed probability – the uniform encoder at 0.8 dropout probability correctly expands both `for i in rng(10):` and `for i in rg(10):`.

But we are still missing something in our model of how humans would choose to omit characters when using our auto-completion framework. Certain characters are more likely to be dropped in accordance with how redundant they are semantically – we approximate this using our *Frequency* encoder.

Consider the Python line `parser.add_option("-v", "-version", help="use a specific zc.buildout version")`, and the encoding `per._opn("-v", "-vern", help="use a specific zc.bdout vern")`. *Uniform* incorrectly decodes to `per._open("-v", "-version", help="user a specific zc.bidout version")`, while *Frequency* decoder expands it correctly. The common 5-grams that were identified in the original string were 'parse', '.add_', 'ption', 'rsion', 'build', and 'rsion' again. The uniform case does not reconstruct 'zc.buildout' from 'zc.bdout' because it treats the deletion of 'u', 'i', and 'l' as independent events, and therefore underestimates the true probability of dropping all three.

The frequency-based decoder sees 10, and guesses that it really comes from the sequence 10000. Even worse, in the beam search the number is further expanded because it thinks that 00 comes from 00000.

This could be due to the decoder learning that it needs to expand the input string, and will prefer longer explanations as a result. The frequency-based decoder's strength is building translations using fairly non-local information; however, the price is that it tends to violate Occam's Razor.

We would like to evaluate these models, but accuracies of different models each using their own encoding schemes are somewhat incomparable. For example, a rule-based encoding is much easier to expand than a uniform encoding. A uniform decoder would have to be very powerful to score as well on uniform encodings as rule based decoder performs on its rule based encodings. We also need to understand which of the encoding schemes is closest to what programmers would actually type when using these auto-completion facilities.

As Table 1 shows, *Rules* decodes its own encodings nearly perfectly, but cannot decode uniform or frequency – it is overfitting to its own very simple encodings. We can see the simplicity of the rule-based encodings in the high uniform and frequency based decoding scores.

The uniform decoder performs meagerly on its own encodings, but performs well on rule-based encodings. The uniform decoder always performs at least tolerably on *Frequency*, with a 27% top-5 accuracy. The uniform decoder can be viewed as underfitting its own encoder, partially because the uniform encoder is especially difficult to reliably decode.

Contrast this with the frequency decoder. It performs well on itself and reasonably on rule-based, but it completely fails on uniform encodings. In other words, it is overfitting its own encoder. The

frequency decoder is inflexible – a small perturbation causes it to leave the regime where it can give confident predictions. This is because the frequency encoder is forced to only consider fixed deletions of n-grams in a fixed order, given by relative n-gram frequencies in the dataset².

We can find a middle ground between the under and over fitting of uniform and frequency based encoders respectively by introducing noise into the frequency encoding, as we did in our *NoisyFrequency* encoder.

To really get a feel for the efficacies of these learned decoders we will need to see examples. Consider translations of `for i in rng(10):` in Table 3.

Rules	'for i in range(10):'	'for i in range(10):ul::'
Uniform	'for i in range(10):'	'for i in range(100):'
Frequency	'for i in rading(10):'	'for i in rng(10000):'
NoisyFrequency	'for i in range(10):'	'for i in range(100):'

Table 3: Top-2 Translations for the input `for i in rng(10):`.

Table 4 contains translations of `r i i r10):`.

Rules	'r i i r10):'	'r i i 10):'
Uniform	'for i in r10):'	'for i in r100000):'
Frequency	'rum, i == r10):'	'rum, i += r10):'
NoisyFrequency	'for i in range(10):'	'for i in r10(self):'

Table 4: Top-2 Translations for `r i i r10):`. R = Rules-Based encoder, U = Uniform Encoder at .8 compression, F = Frequency Encoder dropping 5-grams to compression of .8. NF = Noisy Frequency Encoder dropping 5-grams to compression of .8. All models are trained with lines of Python code.

The decoders that do the right sort of expansions in practice are uniform and noisy frequency, even though noisy frequency didn't come out amazingly well in Table 1 – for example, it has the worst top-1 accuracy of on rule-based encodings. In the last example, the n-gram training helps the NoisyFrequencyEncoder recover the full word 'range' from just 'r', and it actually *completely recovers the target string from r i i r10):*. The extra ability of uniform and Noisy Frequency will come out clearer when evaluating whether the correct answer is in the top five predictions. Indeed, in top-5 accuracy, *NoisyFrequency* performs decently on each type of encoding, and once again beats uniform in self-accuracy. It is no longer the worst at decoding rule-based encodings. That said, if anything, top-5 accuracy further affirms the accuracy of the plain uniform encoder.

We conclude our analysis by showing the uniform and noisy frequency decoders at some of their best and worst examples in Table 5. The best examples picked maximize the number of characters that the decoder had to insert, and are cases in which the decoder was successful. In contrast, the worst examples minimize the number of characters that needed to be inserted, and show cases of failures.

6 Conclusion and Future Work

In this project, we built a keyword-based auto-complete system that is able to make highly non-trivial inferences, as demonstrated in complex examples. We believe that the current project would be usable in a text editor by displaying our few best completion predictions. If the authors were using this project for auto-completion in a real-world setting, we would use either *NoisyFrequency 0.8* or *Uniform 0.7*.

When looking at the output, Uniform 0.7 is able to get common-sense answers more often than Uniform 0.8, and it seems to learn more of a language model when training to predict so many missing characters. Top-1 and top-5 accuracies are still missing an important piece of the evaluation for this project. We also considered a way to measure distance from desired string and use that as a metric, but these results did not reveal any other clear insights that were not evident from accuracies.

²With the caveat in choosing which of multiple equivalently prevalent n-grams to choose

Decoder	Kind	Translation
Uniform 0.8	Best	<pre>BASEPTH = o.ath.dirnamo.patabpath__fil__)</pre> <p style="text-align: center;">↓</p> <pre>BASE_PATH = os.path.dirname(os.path.abspath(__file__))</pre>
Noisy Freq 0.8	Best	<pre>fotcreator = os.os.pth.dirnme(___, '..', fontcreator.py'</pre> <p style="text-align: center;">↓</p> <pre>fontcreator = os.path.join(os.path.dirname(__file__), '..', 'fontcreator.py'</pre>
Uniform 0.8	Worst	<pre>tod_ = smm.tod → tod_ = smm.tod</pre> <p style="text-align: center;">Instead Of: <code>todo_ = smm.todo</code></p>
Noisy Freq 0.8	Worst	<pre>coord = [] → oord = []</pre> <p style="text-align: center;">Instead Of: <code>ood_list = []</code></p>

Table 5: Uniform and Noisy Frequency decoders, at their best and worst performances

We might be able to quantify the performance of noisy frequency 0.8 and uniform 0.7 by using a language model on the translations, which remains as future work.

We also tried uniform 0.6 and a uniform curriculum where we decrease the probability of keeping characters during training. These led to some reasonable but not outstanding expansions. We believe that it would be difficult to do considerably better than uniform 0.7 and noisy frequency 0.8 for translation through further ad-hoc methods. However, we would not be surprised if successfully learning a neural encoder-decoder system end-to-end, as done in [?], would lead to significantly better performance. Trying other variants of that system that might work better with character-level input is also left to future investigation.